

Siddharth Deepak

Leander, TX • (737) 368-2120 • sid.deepak2023@gmail.com • [linkedin.com/in/siddharthdeepak](https://www.linkedin.com/in/siddharthdeepak) • github.com/siddeepak2023 • portfolio: siddharthdeepak.bizfalconai.com

EDUCATION

B.S. in Business Analytics and AI; Concentration in Marketing and AI

Anticipated Graduation: Dec 2027

The University of Texas at Dallas, Naveen Jindal School of Management, Richardson, TX

Activities: AI in Business Club • UTD OCCM • UTD Focus

RELEVANT COURSEWORK

Data Science for Business Applications, AI in Business, Business Analytics, Marketing Analytics, Foundations of Business Intelligence, Database Analytics, Object-Oriented Programming in Python, Quantitative Business Analysis

TECHNICAL SKILLS

Languages:	Python, SQL, TypeScript, HTML/JS/CSS (familiar with R)	Libraries:	Pandas, NumPy, scikit-learn, statsmodels, HistGradientBoosting, Optuna
Tools:	Tableau, Excel, Git, SQLite, Three.js, Chart.js, Firebase	Concepts:	ML Modeling, Model Validation & Calibration, A/B Testing, Statistical Hypothesis Testing, Feature Ablation, Time-Series Forecasting, EDA, Data Cleaning

PROFESSIONAL EXPERIENCE

BizFalcon AI bizfalconai.com, Founder & Analytics Lead

May 2026 – Present

- Launched two analytics products: **ProfitFalcon** for e-commerce brands, covering contribution margin, LTV/CAC, discount leakage and stockout risk, and **ClientFalcon** for marketing agencies, multi-client reporting built on an explainable account-health score. Both run on standard exports from Shopify, Meta, Google Ads, GA4 and Klaviyo
- Engineered the shared analysis engine, **Talon**, which runs Python across a client's full dataset rather than a sample. Authored **9 industry metric packs** and **11 guided statistical workflows** covering hypothesis testing, regression, cohort and time-series analysis, exposing the formula behind every figure for audit
- Designed a recommendation layer that names the next action, such as cutting a budget or reordering a SKU, with **projected dollar impact** and required human approval, then scored each decision against its realized result to make recommendation quality measurable
- Recruited and led a **3-person team** across engineering and operations while personally building and operating the full stack in React, TypeScript, Firebase, in-browser Python and Stripe, maintaining **2,400+ automated tests** under continuous integration

YMCA Central Texas, Summer Camp Counselor

May 2021 – Aug 2025

- Led groups of 15 to 25 campers across four summers and trained incoming counselors through a structured shadowing program
- Earned Counselor of the Week five times and Rookie Counselor of the Year for leadership and camper experience

PROJECTS

Retail Customer Intelligence Platform

[Live Dashboard](#), [GitHub](#)

Python, scikit-learn, Pandas, RFM segmentation, market-basket analysis

- Analyzed **802,712 real transactions** across 41 countries and £17.45M in revenue, segmenting customers into four tiers with K-means RFM and training a random-forest churn model that scored **0.796 ROC-AUC out of sample against a 55.3% base rate**
- Sized the retention campaign with a **precision-at-k curve** showing the top 100 contacts are 90% true churners, and a cost model identifying contact capacity rather than budget as the binding constraint. Generated Apriori market-basket rules with a top pair lift of 24x for cross-sell

NBA Game Outcome Prediction Model

[Live Dashboard](#), [GitHub](#)

Python, HistGradientBoosting, Optuna, SQLite, nba_api

- Trained a gradient-boosted classifier on **10,624 real NBA games** across 9 seasons and validated it on the fully held-out 2024-25 season, reaching **66.4% accuracy against a 54.4% always-home baseline** with calibrated probabilities, TimeSeriesSplit, and 2020 bubble games down-weighted
- Ran a **feature ablation** showing Elo differential alone reaches 0.710 AUC, so the remaining 70 features added **0.015, less than the 0.021 spread across folds**. Re-tested on **1,230 unseen games** from 2025-26, where the model matched a closed-form Elo on ranking and beat it only on calibration, establishing probability quality rather than ranking as its value

Olist E-Commerce Revenue & Sales Analysis

[Live Dashboard](#)

Python, SQL, Pandas, Tableau, Excel

- Analyzed **96,478 e-commerce orders** across 9 source tables, joining them into a star schema and querying in SQL with CTEs and window functions to rank Health and Beauty as the top revenue category at R\$1.4M of R\$15.4M
- Modeled review scores with a logistic regression controlling for state, category and freight, finding each day of delivery delay raised the odds of a low review by **8.5%, or 1.78 times over a one-week slip** (OR 1.085/day, 95% CI 1.083-1.088, n=95,824). Published as a live Tableau dashboard