

SIDDHARTH DEEPAK

Leander, TX • (737) 368-2120 • sid.deepak2023@gmail.com
[linkedin.com/in/siddharthdeepak](https://www.linkedin.com/in/siddharthdeepak) • github.com/siddeepak2023 • siddharthdeepak.bizfalconai.com

EDUCATION

The University of Texas at Dallas — Naveen Jindal School of Management, Richardson, TX

Expected Dec 2027

B.S. in Business Analytics and AI | Concentration in Marketing and AI

Relevant Coursework: Data Science for Business Applications, AI in Business, Marketing Analytics, Database Analytics, Business Intelligence, Quantitative Business Analysis

Activities: AI in Business Club

TECHNICAL SKILLS

Languages: Python, SQL, TypeScript (coursework in R, C++)

Tools & Libraries: Pandas, NumPy, scikit-learn, statsmodels, React, Tableau, Excel, Git, SQLite, Firebase

Methods: Machine learning, model validation, A/B testing, hypothesis testing, logistic regression, customer segmentation, data cleaning

PROFESSIONAL EXPERIENCE

BizFalcon AI — Founder & Analytics Lead

May 2026 – Present

Analytics products for e-commerce brands and marketing agencies | bizfalconai.com

- **ProfitFalcon (business analytics)** computes unit economics for e-commerce brands out of Shopify, Amazon and WooCommerce order exports: contribution margin by SKU, LTV:CAC by acquisition cohort, discount leakage by promo code, and reorder timing from sell-through rates.
- **ClientFalcon (marketing analytics)** reports across an agency's whole client book from Meta, Google Ads and GA4 exports. It ranks every account on an explainable health score built from ROAS against that client's breakeven, the trend since the last report, and how long the agency has gone without sending one.
- **BizFalcon Analyst (free tier)** is the engine under both, and the way in. It runs over a user's entire dataset instead of a sample, and shows the formula behind every figure it prints.
- Built the recommendation layer on top: it prices each proposed action in dollars, holds it for human approval, then scores what actually happened against what it projected, so the quality of the advice is measurable rather than asserted.
- Engineered the analysis engine and the full stack in Python, React/TypeScript, Firebase and Stripe, with 3,000+ automated tests under continuous integration. Sole engineer through August 2026; briefly led a three-person team.

YMCA Central Texas — Summer Camp Counselor

May 2021 – Aug 2025

- Supervised groups of 15–25 campers across four summers and trained incoming counselors. Counselor of the Week (5×) and Rookie Counselor of the Year.

PROJECTS

Retail Customer Intelligence Platform | Python, scikit-learn, Pandas

[Live Dashboard](#)

- Segmented 802,712 transactions across 41 countries (£17.45M) into four value tiers using K-means RFM analysis, then built a churn model reaching 90% precision among top-ranked contacts.
- Sized the retention campaign with a cost model that identified outreach capacity, not budget, as the binding constraint; market-basket rules (24× top-pair lift) drove cross-sell targeting.

Olist E-Commerce Revenue & Delivery Analysis | Python, SQL, Tableau

[Tableau Dashboard](#)

- Modeled 96,478 orders from 9 source files into an 8-table star schema, querying it in SQL with CTEs and window functions.
- Quantified delivery impact on satisfaction: a one-week delay raised the odds of a negative review 1.78× after controlling for state, category, and freight cost. Published as a live Tableau dashboard.

NBA Game Outcome Prediction | Python, scikit-learn, SQLite

[GitHub](#)

- Trained a gradient-boosted classifier on 10,624 games across 9 seasons and validated it on a fully held-out season, reaching 66.4% accuracy against a 54.4% baseline.
- Ran a feature ablation demonstrating that 70 of 71 features contributed less signal than run-to-run variance, redirecting effort toward probability calibration rather than added complexity.

Universe Explorer | Python, Pandas, Astropy

[Live Demo](#)

- Cleaned and merged 10.2M records from four published astronomical catalogues (Gaia DR3, SDSS DR17, 2MRS, NASA Exoplanet Archive): dropped nulls and low-quality measurements, converted parallax to distance, reconciled coordinate systems, and deduplicated across catalogues into the 4.6M-object dataset the live site serves.